

Part E: Mortgage Survival Analysis

Competing Risks · Time-Varying Cox · Neural Models · Scenario Analysis

Yueqi Lin

MTH 9877 — Interest Rate and Credit Models
Baruch College, Spring 2026



May 13, 2026

Part E(i)

Competing Risks: Prepayment vs. Default

Kaplan–Meier vs. Aalen–Johansen · Cause-Specific Cox · Fine–Gray

Prepaid and default are competing risks.



Kaplan–Meier vs. Aalen–Johansen.

Kaplan–Meier

$$\hat{F}_{\text{KM}}^{(k)}(t) = 1 - \prod_{t_j \leq t} \left(1 - \frac{d_j^{(k)}}{n_j^{\text{KM}}} \right)$$

$d_j^{(k)}$ cause- k exits at t_j
 n_j^{KM} **risk set excluding competing-event loans**

Aalen–Johansen


$$\hat{F}_k(t) = \sum_{t_j \leq t} \hat{S}(t_j^-) \cdot \frac{d_{kj}}{n_j}$$

d_{kj} cause- k exits at t_j
 n_j risk set including all exits
 $\hat{S}(t_j^-)$ **overall survival just before t_j**

KM: competing exits censored $\Rightarrow$ denominator too small $\Rightarrow$ each increment inflated.

AJ: $\hat{S}(t_j^-)$ shrinks with every exit $\Rightarrow \hat{F}^{(\text{pre})} + \hat{F}^{(\text{def})} + \hat{S} = 1$ always.

Aalen–Johansen CIF: coherent probability partition.



Cumulative Incidence Function


$$F_k(t) = P(\text{exit by } t, \text{ cause } k)$$

Aalen–Johansen CIF partition

	Prepaid	Default	Active
5 yr	61.4%	1.1%	37.6%
10 yr	83.3%	2.2%	14.5%
20 yr	91.8%	3.1%	5.1%

KM overstates prepayment by up to 2.2 pp. over 20 years.

KM over-estimates prepayment CIF by treating defaults as random censoring









KM overstatement of prepayment probability

Horizon	$\hat{F}_{KM}$	$\hat{F}_{AJ}$	Bias
5 yr	61.9%	61.4%	+0.45 pp
10 yr	84.6%	83.3%	+1.31 pp
20 yr	93.4%	91.8%	+2.23 pp

FICO and LTV separate default, not prepayment.

Stratified Competing-Risk CIFs [5% sample]



Cause-specific Cox model.

Cause-Specific Cox

$$h_k(t | \mathbf{x}) = h_{0k}(t) \exp(\boldsymbol{\beta}_k^\top \mathbf{x}), \quad k \in \{\text{pre}, \text{def}\}$$

$$\ell(\boldsymbol{\beta}_k) = \sum_{i: E_i=k} \left[\boldsymbol{\beta}_k^\top \mathbf{x}_i - \log \sum_{j \in \mathcal{R}(t_i)} e^{\boldsymbol{\beta}_k^\top \mathbf{x}_j} \right]$$

$\mathcal{R}(t_i)$: all loans still active just before t_i .

Competing-cause exits are **censored** — they leave $\mathcal{R}(t)$ without contributing to ℓ .

Additional features


Max Delinquency — months past due over loan life

Ever 90+ DQ — binary flag for severe delinquency


Opposite covariate effects by cause.

Prepayment and default respond to opposite covariate effects

Cause-Specific Cox — Prepayment



Cause-Specific Cox — Default



Four covariates flip sign across causes.

Naive Cox vs. cause-specific coefficients

Covariate	Naive $\hat{\beta}$	Prepay $\hat{\beta}$	Default $\hat{\beta}$	
Rate incentive	+0.253	+0.256	+0.021	(pure prepay)
UPB	+0.161	+0.161	-0.010	(pure prepay)
FICO	+0.016	+0.010	-0.047	opposite
LTV	-0.007	-0.011	+0.041	opposite
Max Delinquency	-0.123	-0.245	+0.298	opposite
Ever 90+ DQ	-0.035	-0.071	+0.270	opposite

Each extra month delinquent raises default hazard by 35% and cuts prepayment hazard by 22% — yet the naive coefficient (-0.123) conceals both, collapsing a $+0.543$ spread to near zero.

$$e^{+0.298} - 1 = +35\% \text{ (default); } 1 - e^{-0.245} = +22\% \text{ (prepay reduction); } \text{spread} = 0.298 - (-0.245) = 0.543.$$

Key Takeaways: $E(i)$

- ① Default and prepayment are competing risks — a naive Cox model conflates both.
- ② **Kaplan–Meier overstates prepayment by up to +2.2 pp over 20 years; use Aalen–Johansen.**
- ③ Vintage era dominates: 2007 cohort defaults at $6\times$ the long-run average.
- ④ FICO drives default; rate incentive drives prepayment — separate channels.

Part E(ii)

Time-Varying Covariates in Mortgage Survival

Andersen–Gill Cox Extension · Live Rate Incentive · Landmark C-index

Static Cox vs. Andersen–Gill.

Standard Cox (11 features)

One row per loan; all covariates fixed at origination.
Cannot track rate cycles, burnout, or equity build after origination.

Andersen–Gill (13 features)

One row per loan-month; $x_i(t)$ updated each period.
Extends the same static base with the four time-varying features, tracking the full path of rate incentive and equity. 100K loans $\times$ ~ 47 months = 4.68M intervals.

Time-varying features — updated every month

Feature	What it captures
rate_incentive	$r_{\text{orig}} - r_{\text{mkt}}(t)$: live option value to refinance; tracks rate cycles and burnout
ELTV	equity from amortisation and HPI gains; high ELTV suppresses refi
unemployment	macro labour cycle; proxies Fed easing in high-unemp periods
hpi_yoy	home-price appreciation unlocking cash-out refinancing

*unemployment and hpi_yoy also in Standard Cox as origination-month snapshots.

The Andersen–Gill model.

Hazard function (identical in form to standard Cox)

$$h_i(t) = h_0(t) \exp(\beta^\top x_i(t))$$

$x_i(t)$ is the **current** covariate vector, updated each month. β is a single coefficient vector estimated across all loan-months.



Log partial likelihood (maximised to obtain $\hat{\beta}$)

$$\ell(\beta) = \sum_{i: \delta_i=1} \left[\underbrace{\beta^\top x_i(t_i)}_{\text{log-score of the prepaying loan}} - \log \underbrace{\sum_{j \in \mathcal{R}(t_i)} \exp(\beta^\top x_j(t_i))}_{\text{sum over all loan-months active at } t_i} \right]$$

At each event: *given a prepayment at t_i , what is the probability it was loan i rather than any other active loan?*
Maximising ℓ estimates $\hat{\beta}$ without specifying $h_0(t)$.

Time-varying feature trajectories by vintage.


Static Cox uses only month-0 values — Andersen-Gill tracks these every month



Static Cox uses only month-0 values — Andersen-Gill tracks these every month.


Static vs. TV Cox coefficient comparison.

Feature coefficients: TV Cox (current monthly) vs Static Cox (origination snapshot)
Shared static features (Orig. rate, UPB) shown in both panels



- `rate_incentive`: dominant positive predictor in TV Cox; near zero at origination, absent from Static Cox
- Macro signals (unemployment, HPI) flip sign — vintage-era snapshots capture a different regime than monthly values

Landmark C-index: TV Cox leads to $T=60$, then reverses.



Landmark C-index

$$C(t^*) = P(\hat{h}_i(t^*) > \hat{h}_j(t^*) \mid T_i \leq t^* < T_j)$$

Horizon	Static	TV	Δ
$T = 12$	0.609	0.681	+0.072
$T = 24$	0.627	0.687	+0.060
$T = 36$	0.630	0.675	+0.045
$T = 60$	0.625	0.646	+0.021
$T = 84$	0.623	0.605	-0.018
$T = 120$	0.620	0.583	-0.037

Key Takeaways: E(ii)

- ① Monthly-updated covariates add real information (+1,996 log-likelihood gain).
- ② **TV Cox leads at short/medium horizons: +0.072 at $T=12$, +0.021 at $T=60$.**
- ③ **TV Cox degrades beyond $T \approx 70$ months — use a landmark super-model.**
- ④ Rate incentive is the key driver; macro coefficients require careful inspection.

Part E(iii)

Neural Survival Models: LongitudinalDeepHit

Transformer Encoder · Competing-Risk CIFs · Gradient Sensitivity

Why go beyond Cox?

Cox / Andersen-Gill limitations

- Proportional hazards: covariate effect is **constant over time**
- Single scalar score — cannot model the **S-curve shape** of prepayment
- Competing risks require separate model fits; no shared representation
- Cannot exploit the **sequence structure** of monthly covariate history

What a sequence model can do

- Read the full monthly path $[x_1, \dots, x_T]$ — **burnout is encoded**
- Produce nonlinear, time-varying hazard rates $\lambda_k(t \mid x)$
- Model prepayment and default **jointly** via shared encoder
- Attention weights identify **which months** drive the prediction

Alternatives rejected

- **LSTM** — vanishing gradients at $T=120$; attends only to adjacent months; hard to interpret
- **Single-context head** — no time index in head; near-constant hazard curve; gradient flows only through last hidden state
- **Two separate models** — violates independent censoring (defaulted loans are riskiest borrowers, not random exits); loses shared macro signal

Why this combination

- **Causal Transformer** — reaches any past month in one attention step; gradient flows through every position; weights are interpretable
- **Per-timestep head** — $\lambda(t)=\text{head}(h_t)$: hazard conditioned on full causal history $x_1 \dots x_t$; every month trained independently
- **Joint encoder** — shared representation captures that the same macro shock raises default and suppresses prepayment; chain product enforces $\text{CIF}_1 + \text{CIF}_2 + S = 1$

LongitudinalDeepHit architecture.

Architecture

Input	$(B, T_{\max}=120, D=13)$ padded sequences
Proj.	Linear $D \rightarrow d_{\text{model}}=64$
PE	Sinusoidal positional encoding on loan age
Enc.	$\times 2$ pre-norm layers, 4 heads, causal mask
Head 1	FC $64 \rightarrow 64 \rightarrow 1$ per h_t , Sigmoid $\Rightarrow \lambda_1(t)$
Head 2	FC $64 \rightarrow 64 \rightarrow 1$ per h_t , Sigmoid $\Rightarrow \lambda_2(t)$
CIFs	Competing-risk chain product

Competing-risk CIFs

$$\begin{aligned}\text{CIF}_k(t) &= \sum_{s=1}^t \lambda_k(s) \cdot S(s-1) \\ S(t) &= \prod_{s=1}^t (1 - \lambda_1(s))(1 - \lambda_2(s))\end{aligned}$$

Guarantees $\text{CIF}_1 + \text{CIF}_2 + S = 1$ at every t .

Training objective

$$\mathcal{L} = \mathcal{L}_{\text{NLL}} + \alpha \mathcal{L}_{\text{rank}}, \quad \alpha = 0.2$$

NLL: discrete-time cause-specific likelihood

Ranking: soft pairwise concordance (DeepHit-style)

Params	76,290
Data	101K loans, 4.5M rows
Epochs	41 (best ep. 29)
LR	3×10^{-4} , cosine
Loss	11.30 $\rightarrow$ 11.28
Rank	1.97 $\rightarrow$ 1.84

Data treatment


Each loan $\rightarrow$ 13-feature monthly sequence, zero-padded to $T_{\max}=120$ after exit.

Labels: prepayment = 1, default = 2, censored = 0.

Causal mask; 80/20 train-test split.

Training dynamics: Negative Log Likelihood collapse and ranking loss.



LongitudinalDeepHit — Training Curves



- **NLL freezes:** 1.6% event rate means predicting near-zero hazard for everyone already minimises the loss — rare-event gradients too weak to improve further.
- **Cox immune:** partial likelihood only ranks loans at event times, never predicts absolute hazard — no trivial solution under class imbalance.
- **AUC preserved:** ranking loss descends steadily, ordering loans by relative risk despite collapsed hazard values.

Gradient sensitivity by cause.

Gradient Sensitivity Analysis — LongitudinalDeepHit




Red = time-varying features. Blue = static features.

Dynamic features rank above static for prepayment;

FICO and unemployment lead for default; consistent with cause-specific Cox in E(i).

Feature importance over the loan lifecycle.



rate_incentive and Mortgage rate dominates, especially in early months;

Model performance vs. DeepCox.

Prepayment CIF₁ — AUC at fixed horizons

Model	T=12	T=24	T=36	T=60
Part D DeepCox [†]	0.684	0.701	0.721	0.729
DeepHit (20K) [‡]	0.641	0.622	0.572	0.515

Default CIF₂ — AUC at fixed horizons

Model	T=12	T=24	T=36	T=60
Part D DeepCox [†]	0.684	0.701	0.721	0.729
DeepHit (20K) [‡]	0.687	0.708	0.707	0.726

[†] 597K canonical [‡] 20K hold-out green = better

Prepayment: DeepCox leads at all horizons (NLL collapse hurts long-horizon calibration).

Default: DeepHit leads at T=12 and T=24 with **6×** fewer training loans.

Key Takeaways: E(iii)

- 1 Default AUC matches DeepCox with $6\times$ fewer training loans.
- 2 Gradient sensitivity replicates Cox findings: rate incentive drives prepayment, FICO drives default.
- 3 Feature importance shifts over the loan lifecycle — unique to sequence models.
- 4 **Limitation: defaults are rare (1.6%), so the model learns to predict near-zero hazard for everyone — AUC is preserved but absolute probabilities are unreliable; fix via focal-loss upweighting.**

Part E(iv)

Scenario Analysis: MBS Prepayment under Rate Shocks

TV Cox Rate Channel · Survival Curves · Monte Carlo MBS Pricing

TV Cox: direct rate channel for scenario analysis.

Why TV Cox (E(ii)) is the primary model

$\text{rate_incentive} = r_{\text{orig}} - r_{\text{market}}(t)$ is a **direct time-varying covariate**.

A Δr bp market rate shock immediately shifts the incentive by $-\Delta r/100$ pp, giving an analytic hazard response:


$$\Delta \log h(t) = \hat{\beta}_{\text{ri}} \times \frac{-\Delta r}{100}, \quad \hat{\beta}_{\text{ri}} = 0.041 \text{ (HR} = 1.042 \text{ per 100 bp)}$$

The calibrated baseline hazard $\hat{H}_0(t)$ produces realistic $S(t)$, CPR, WAL, and price — no approximation needed.

Representative loan: \$300K, 5.0% coupon, baseline market rate 5.3% (median, 1999–2025).

Scenarios: ± 100 , ± 200 , ± 300 bp — covering the BCBS standard stress range.


Shock sensitivity: TV Cox matches the analytic formula.



Each 100 bp rate cut increases the prepayment hazard by 4.2%; a 300 bp cut raises it by +0.123 in log terms.
The model response matches the closed-form prediction exactly — no simulation needed, fully auditable.

Survival curves under seven rate scenarios.



Prepayment Survival Curves Under Interest Rate Scenarios (TV Cox)
(Representative loan: \$300K, 6.5% coupon, baseline market rate 7.0%)



Scenario	10-yr Survival	Median Prepay (mo.)
–300 bp	10.2%	41
–200 bp	11.2%	43
–100 bp	12.2%	45
Baseline	13.3%	46
+100 bp	14.4%	48
+200 bp	15.6%	50
+300 bp	16.8%	51

MBS: Weighted Average Life, MBS price, and negative convexity.

MBS Cash Flow Projection — TV Cox Monte Carlo ($\sigma=25$ bp, 200 paths, 10-yr horizon)
5.0% coupon, \$300K loan, baseline market rate 5.3%



Scenario	WAL (yr)	Price (\$)
−300 bp	3.25	108.32
−200 bp	3.34	105.26
−100 bp	3.43	102.17
Baseline	3.53	99.07
+100 bp	3.63	95.96
+200 bp	3.73	92.86
+300 bp	3.83	89.77

Negative convexity: when rates fall, borrowers prepay and return principal *at par*, capping price upside. Price gain (−300 bp: \$9.25) < loss (+300 bp: \$9.30).

MC: 200 paths, $\sigma=25$ bp.

Key Takeaways: E(iv)

- 1 **TV Cox has a linear rate channel:** every 100 bp rate cut adds the same +4.2% to prepayment hazard — equal steps in rate produce equal steps in response, making results predictable and fully explainable.
- 2 **Negative convexity confirmed:** when rates fall, borrowers prepay at par, capping price upside (\$9.25 gain < \$9.30 loss at ± 300 bp).

Thank you

Questions?

Andersen & Gill (1982) *Ann. Stat.* · Fine & Gray (1999) *JASA* · Lee et al. (2018) *DeepHit* · Sadhwani et al. (2021) *J. Fin. Econometrics* · BCBS (2019)